

Scaling computations using parallel computing

Przemysław Szufel
<https://szufel.pl/>

Preparation of this workshop has been supported by the Polish National Agency for Academic Exchange under the Strategic Partnerships programme, grant number BPI/PST/2021/1/00069/U/00001.



POLISH NATIONAL AGENCY
FOR ACADEMIC EXCHANGE

Parallelization options in programming languages

- Single instruction, multiple data (SIMD)
- **Green-threads** (co-routines) only where I/O is bottleneck
- **Multi-threading**
 - Language
 - Libraries
- **Multi-processing** – many julia processes at the same time
 - single machine
 - distributed (cluster)
 - distributed (cluster) via external tools
- GPU computing

Comparison of parallelism types

Threading

- Single process (cheap)
- Shared memory
- Number of threads running simultaneously limited by number of processors
- Possible issues with locking and false sharing

Multiprocessing

- Multiple processes
- Separate memory
- Number of processes running simultaneously limited by cluster size
- Possible issues if inter-process communication is needed

Distributed computing

- Multiple hosts
- Multiple processes (single- or multi-threaded)

Julia command line option for parallelism

multi-threading

`-t, --threads {N|auto}`

Enable N threads; "auto" currently sets N to the number of local CPU threads but this might change in the future

multi-processing

`-p, --procs {N|auto}`

Integer value N launches N additional local worker processes


"auto" launches as many workers as the number of local CPU threads (logical cores)

distributed computing


`--machine-file <file>`

Run processes on hosts listed in <file>

Threading



Single Process P with single thread



Single Process P with three threads

Simple example – threading

Single threaded

```
function ssum(x)
    r, c = size(x)
    y = zeros(c)
    for i in 1:c
        for j in 1:r
            @inbounds y[i] += x[j, i]
        end
    end
    y
end
```

Multithreading

```
function tsum(x)
    r, c = size(x)
    y = zeros(c)
    Threads.@threads for i in 1:c
        for j in 1:r
            @inbounds y[i] += x[j, i]
        end
    end
    y
end
```

Sum: output

threads: 1

0.721701 seconds (2 allocations: 156.328 KiB)

0.858513 seconds (2 allocations: 156.328 KiB)

0.730469 seconds (60.27 k allocations: 3.834 MiB, 8.09% compilation time)

0.914933 seconds (8 allocations: 156.938 KiB)

} delta is compilation time

threads: 2

0.750043 seconds (2 allocations: 156.328 KiB)

0.882158 seconds (2 allocations: 156.328 KiB)

0.415769 seconds (60.29 k allocations: 3.835 MiB, 14.25% compilation time)

0.341232 seconds (14 allocations: 157.469 KiB)

threads: 4

0.709045 seconds (2 allocations: 156.328 KiB)

0.690816 seconds (2 allocations: 156.328 KiB)

0.261915 seconds (60.32 k allocations: 3.838 MiB, 23.98% compilation time)

0.224480 seconds (24 allocations: 158.484 KiB)

Threading: synchronization

Increment x 10^6 times using threads:

- Atomic operations
- SpinLock (busy waiting)
- Mutex (OS provided lock)

Locking: output

1 thread

Row	f	i	value	timems
1	f_bad	1	10000000	65.416
2	f_bad	2	10000000	47.856
3	f_atomic	1	10000000	12.144
4	f_atomic	2	10000000	5.9725
5	f_spin	1	10000000	116.92
6	f_spin	2	10000000	107.22
7	f_reentrant	1	10000000	175.71
8	f_reentrant	2	10000000	180.96

4 threads

Row	f	i	value	timems
1	f_bad	1	2500984	32.145
2	f_bad	2	2501422	30.683
3	f_atomic	1	10000000	26.352
4	f_atomic	2	10000000	18.803
5	f_spin	1	10000000	186.86
6	f_spin	2	10000000	185.11
7	f_reentrant	1	10000000	998.39
8	f_reentrant	2	10000000	1003.6

Green threading

- In computer programming, green threads are threads that are scheduled by a runtime library or virtual machine (VM) instead of natively by the underlying operating system. Green threads **emulate** multithreaded environments without relying on any native OS capabilities, and they are managed in user space instead of kernel space, enabling them to work in environments that do not have native thread support.

https://en.wikipedia.org/wiki/Green_threads

Green threads in Julia

```
julia> @time sleep(2)
```

```
2.004142 seconds (11 allocations: 320 bytes)
```

```
julia> @time @async sleep(2)
```

```
0.000128 seconds (27 allocations: 2.250 KiB)
```

```
Task (runnable) @0x0000000012c22e10
```

Green threads - sample

```
function dojob(i)
    val = rand()
    sleep(val)    # this could be external computations with I/O
    i, val
end

result = Vector{Tuple{Int,Float64}}(undef, 8)

@sync for i=1:8
    @async x[i] = dojob(i)
end
```

Example – multiprocessing

using Distributed

```
function s_rand()  
    n = 10^4  
    x = 0.0  
    for i in 1:n  
        x += sum(rand(10^4))  
    end  
    x / n  
end
```

```
@time s_rand()  
@time s_rand()
```

```
function p_rand()  
    n = 10^4  
    x = @distributed (+) for i in 1:n  
        sum(rand(10^4))  
    end  
    x / n  
end
```

```
@time p_rand()  
@time p_rand()
```

Rand: output

0.685125 seconds (20.01 k allocations: 763.722 MiB, 16.21% gc time)

0.642642 seconds (20.00 k allocations: 763.702 MiB, 15.37% gc time)

0.927246 seconds (356.68 k allocations: 783.715 MiB, 11.54% gc time, 26.08% compilation time)

0.471032 seconds (20.03 k allocations: 763.704 MiB, 14.47% gc time)

Parallelizing Julia code

- `@distributed`
- `@spawnat`
- `@everywhere`
- `@async`
- `@sync`
- `fetch()`

Typical pattern for distributed computation

```
using Distributed
addprocs(4); # instead -p

@everywhere include("worker_setup.jl")

function init_worker()
    Random.seed!(myid())
end

@sync for wid in workers()
    @async fetch(@spawnat wid init_worker())
end
```


Writing distributed loops


```
data = @distributed (append!) for (i, j) =  
    vec(collect(Iterators.product(1:4, 1:5)))  
    a = rand(1:499)  
    b = rand(1:9)*1000  
    c = calc(a, b)  
    DataFrame(;i,j,a,b,c,procid = myid())  
end
```

BACKGROUND MATERIALS...

Single instruction multiple data
SIMD

SIMD

- Single instruction, multiple data (SIMD) describes computers with multiple processing elements that perform the same operation on multiple data points simultaneously. Such machines exploit data level parallelism, but not concurrency: there are simultaneous (parallel) computations, but only a single process (instruction) at a given moment.



Source: <https://en.wikipedia.org/wiki/SIMD>

Data level parallelism




Image source: <https://en.wikipedia.org/wiki/SIMD>

#1_dot_simd.jl

```
function dot1(x, y)
    s = 0.0
    for i in 1:length(x)
        @inbounds s += x[i]*y[i]
    end
    s
end

function dot2(x, y)
    s = 0.0
    @simd for i in 1:length(x)
        @inbounds s += x[i]*y[i]
    end
    s
end
```

Dot product: output

12.900 μ s (0 allocations: 0 bytes)


1.760 μ s (0 allocations: 0 bytes)

2.4592981764713477e7

2.4592981764713492e7

Configuration issues for
distributed computing
environments

A typical „single server in the cloud” configuration



Connecting to a cloud instance



- Connect to the cloud server

```
$ ssh -i keyfile.pem ubuntu@ec2-18-218-237-1.us-east-2.compute.amazonaws.com
```



- Copy a local file (note the slash type on Windows)

```
$ scp -i keyfile.pem c:\temp\local.txt ubuntu@ec2-18-218-237-1.us-east-2.compute.amazonaws.com:/home/ubuntu/
```



- Copy a local folder (note the slash type on Windows)

```
$ scp -r -i c:\temp\folder ubuntu@ec2-18-218-237-1.us-east-2.compute.amazonaws.com:/home/ubuntu/
```

Notes:

Mac OSX and Linux: run `chmod 600 keyfile.pem` before using the keyfile



Windows: The best SSH/SCP for windows is contained within Git: <https://git-scm.com/download/win>

Typical enviroment

- Cluster controller
 - Machine type: t2.micro (free tier)
 - Number of machines: 1
- Cluster nodes
 - Machine type: c5.large (Spot Fleet)
 - Number of machines: 5
- You should know how to:
 - Basic understanding of HTTP protocol
 - Use console
 - Use SSH client
- All work throughout the workshop will take place in the cloud
(no local software except for an SSH client a web browser is required)

What is a computing cluster


A group of computers working towards achieving a common goal.



A computer within a cluster is called a **node**

Passwordless SSH cluster


- Enables direct execution of commands on other nodes via command:
`ssh user@hostname command`
- Environments such as Gnu R parallel and Julia parallel use this feature to spawn processes on worker nodes
- Control node has the private key while each slave node needs to have the public key in the `~/.ssh/authorized_users` file
- Network connections between nodes should be open
 - Create a SecurityGroup that allows unlimited connections to itself.



Setting up SSH for passwordless SSH

- `ssh-keygen -P "" -t rsa -f ~/.ssh/cluster`
- `printf "\nUser ubuntu\nPubKeyAuthentication yes\nStrictHostKeyChecking no\nIdentityFile ~/.ssh/cluster\n" >> ~/.ssh/config`
- `cat ~/.ssh/cluster.pub >> ~/.ssh/authorized_keys`
- Test configuration:
 - `$ ssh localhost`

Paswordless SSH cluster



Self-referencing SecurityGroup

Name

Group ID

Group Name

VPC ID

Description

<div></div>	sg-ebdb1b80	launch-wizard-40	vpc-c30fe6aa	launch-wizard-40 created 2018-01-25T21:55:12.263+01:...
-------------	-------------	------------------	--------------	---

Security Group: sg-ebdb1b80

Description

Inbound

Outbound

Tags

Edit

Type	Protocol	Port Range	Source	Description
All traffic	All	All	sg-ebdb1b80 (launch-wizard-40)	
SSH	TCP	22	0.0.0.0/0	

Julia cluster specification file and running distributed clusters

```
$ more machinefile_julia
```

```
4*ubuntu@172.31.10.229
```

```
4*ubuntu@172.31.11.44
```

```
4*ubuntu@172.31.0.243
```

```
4*ubuntu@172.31.13.134
```

```
4*ubuntu@172.31.14.219
```


```
$ julia -machine-file machinefile_julia program.jl
```

```
# REQUIRES PASSWORDLESS SSH TO BE CONFIGURED!
```


Use case scenario: Performance of
distributed code in Julia
Cray vs AWS

Schelling (1974) segregation model


- Agents occupy cells of rectangular space
- Two types of agents (e.g. blue and red)
- When not happy with their neighbours randomly relocate




Distributed Schelling segregation model



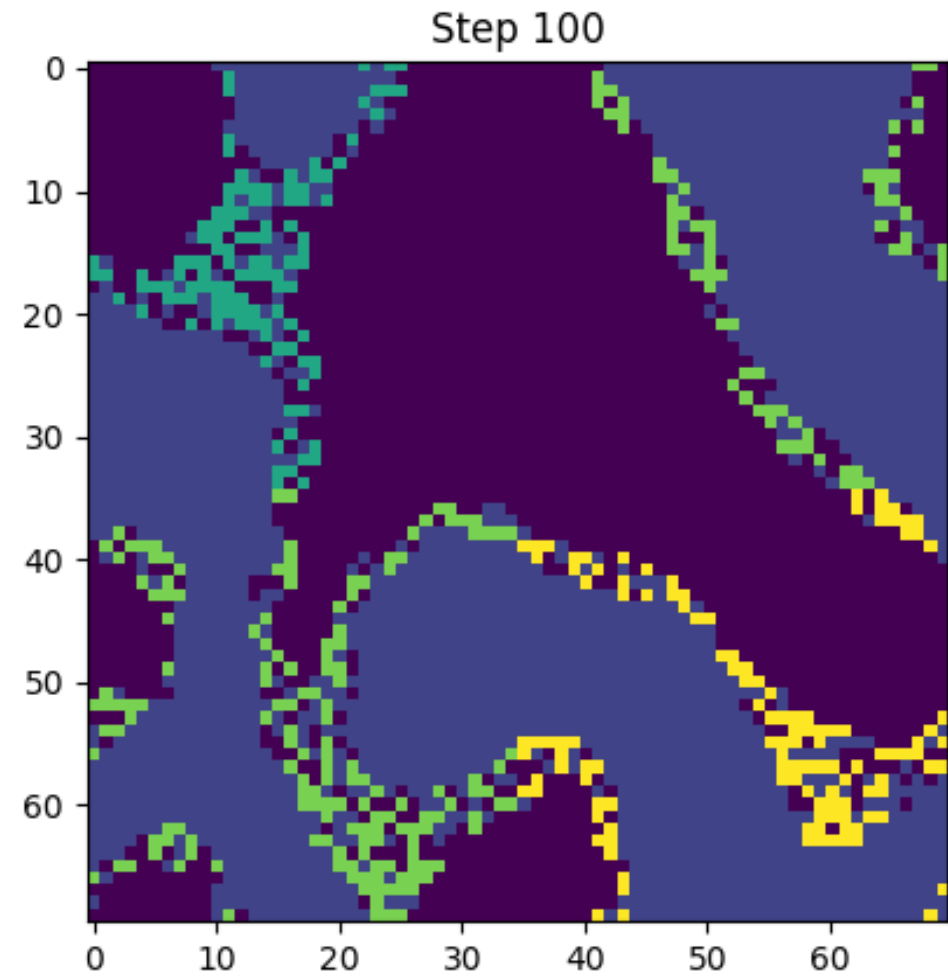

Distributed Schelling segregation model



Distributed simulation architecture



Parallelized Schelling model (2x2)




Parallelizing Julia on Cray with SLURM

```
julia> using ClusterManagers
```

```
julia> addprocs_slurm(200, job_name="my_job",  
    account="GC71-37", time="00:10:00",  
    exename="/lustre/tetyda/home/pszufe/julia/usr/bin/julia")
```


Typical Julia performance pattern



**Time
to
execute
step
(seconds)**

Simulation step

Typical Julia performance pattern




**Time
to
execute
step
(seconds)**

Simulation step

Distributed simulation scalability

**Time
to
execute
step
(seconds)**



Note:

The first step has been removed

Cray vs AWS Spot Fleet

